

TLQLDMR: A Transfer Learning Quantile Regression Model for Wind Power Forecasting under Extreme Weather Conditions

***a, ***a, ***b, ***a and ***a,*

^aSchool of Artificial Intelligence, Southwest University, Chongqing 400715, China

^bBusiness school, Jiangsu open University, Jiangsu 210036, China

ARTICLE INFO

Keywords:

Wind power forecasting
Large margin distribution machine
Quantile Regression
Transfer learning

ABSTRACT

Accurate wind power forecasting under extreme weather conditions is critical for ensuring power grid stability, yet remains fundamentally challenging due to the inherent scarcity of extreme event samples and the substantial distribution shift between normal and anomalous meteorological regimes. While kernel-based methods—particularly Support Vector Regression (SVR) and Large Margin Distribution Machines (LDMR)—have demonstrated strong theoretical foundations and competitive performance in conventional forecasting scenarios, they suffer from three critical limitations when applied to extreme weather prediction: (1) the i.i.d. assumption fails under severe distribution shifts, (2) symmetric error regularization conflicts with the asymmetric nature of quantile-based interval forecasting, and (3) existing formulations lack mechanisms to leverage knowledge transfer from abundant normal weather data to scarce extreme weather samples.

To address these limitations, we propose a novel Transfer Learning Quantile Large Margin Distribution Machine for Regression (TLQLDMR), which extends the classical LDMR framework with three key innovations tailored for cross-domain extreme weather forecasting. First, we introduce an Asymmetric Variance Regularizer that reconciles distributional robustness with quantile estimation by decomposing residual variance into quantile-specific components, thereby preserving the requisite skewness for accurate interval prediction while maintaining margin concentration. Second, we develop a Supervised Transfer Learning framework that integrates Maximum Mean Discrepancy (MMD) for feature-space alignment with a Weighted Pinball Loss mechanism for label-space adaptation, strategically assigning higher weights to scarce target domain samples as corrective anchors. Third, we propose a scalable Nyström-SGD optimization algorithm that reduces computational complexity from $O(N^3)$ to $O(T \cdot N \cdot m)$, enabling practical deployment on large-scale wind farm datasets. Extensive experiments on six wind farms demonstrate that TLQLDMR consistently outperforms both classical kernel-based baselines (SVR, SVQR, LDMR) and state-of-the-art transfer learning methods, achieving 12–18% RMSE reduction and superior 95% interval coverage (94.8%) with optimal reliability-sharpness balance. Comprehensive ablation studies confirm the necessity of each component, establishing TLQLDMR as a principled extension of the LDMR paradigm for transfer learning scenarios.

1. Introduction

The global transition toward renewable energy has positioned wind power as a cornerstone of sustainable electricity generation, with installed capacity projected to exceed 1,000 GW by 2030 (Wang et al., 2019). However, the inherent intermittency and volatility of wind resources pose significant challenges to power grid stability and dispatch optimization. Accurate wind power forecasting is therefore critical for balancing supply and demand, reducing operational costs, and ensuring reliable grid integration (Wang et al., 2019; Zhang & Wang, 2024). While substantial progress has been achieved in short-term forecasting under normal meteorological conditions, prediction performance degrades dramatically during extreme weather events—such as typhoons, cold waves, or heat waves—when wind power fluctuations are most severe and grid stability is most vulnerable (State Grid Renewable Energy, 2022).

Extreme weather events present a dual challenge for data-driven forecasting models. First, such events are inherently rare, resulting in severe data scarcity: extreme samples typically constitute less than 5% of historical records (State Grid Renewable Energy, 2022). Second, the meteorological dynamics during extreme conditions differ fundamentally from normal regimes, inducing substantial distribution shifts in both feature space (e.g., wind speed, temperature) and

*Corresponding author: ***

E-mail address: ***

ORCID(s):

label space (power output). Conventional machine learning models, which assume that training and testing data are independently and identically distributed (i.i.d.), fail catastrophically under such distribution mismatches. For instance, random forest models trained on normal weather data exhibit root mean squared error (RMSE) increases exceeding 40% when applied to extreme weather scenarios (Mustaffa & Sulaiman, 2025; Chen & Shi, 2025).

1.1. Kernel-Based Methods for Wind Power Forecasting: Strengths and Limitations

Among the diverse landscape of forecasting methodologies, kernel-based approaches—particularly Support Vector Regression (SVR) and its variants—have established themselves as a theoretically principled and empirically competitive paradigm. Unlike black-box deep learning models that require massive training data and lack interpretability, kernel methods offer several distinctive advantages: (1) strong generalization guarantees rooted in statistical learning theory (Meena et al., 2025), (2) effective handling of high-dimensional feature spaces through the kernel trick without explicit feature engineering, (3) convex optimization formulations that guarantee global optima, and (4) robust performance on small-to-medium-sized datasets where deep models tend to overfit (Wang & Xu, 2024).

Traditional SVR minimizes the ϵ -insensitive loss to estimate the conditional mean of the target variable, achieving robust point predictions under Gaussian noise assumptions. However, for wind power forecasting under extreme conditions, point estimates are insufficient—grid operators require probabilistic interval forecasts that quantify prediction uncertainty (Hu et al., 2025; Hu & Liu, 2025). To address this need, Support Vector Quantile Regression (SVQR) extends SVR by replacing the symmetric ϵ -loss with the asymmetric pinball loss, enabling estimation of conditional quantiles (e.g., 5th and 95th percentiles for prediction intervals) (Ye et al., 2025b; Wang & Xu, 2024). Recent advances in SVQR include twin formulations for improved computational efficiency (Ye et al., 2023), sparse regularization for feature selection (Wang & Xu, 2024), and nonlinear feature selection mechanisms (Ye et al., 2025b).

Despite these advances, both SVR and SVQR share a fundamental limitation: they optimize only the minimum margin (the distance from the decision boundary to the closest support vectors) while ignoring the statistical distribution of margins across the entire dataset. This focus on worst-case samples can lead to overfitting and poor generalization, particularly when training data contains outliers or label noise. To overcome this limitation, the Large Margin Distribution Machine (LDMR) paradigm was introduced, which explicitly optimizes the first-order and second-order statistics of the margin distribution—maximizing the mean margin while minimizing the margin variance (Zhang & Zhou, 2019; Gupta & Gupta, 2021). By concentrating prediction errors tightly around zero, LDMR achieves superior generalization compared to classical SVR on a wide range of regression tasks (Gupta & Gupta, 2021; Hazarika & Gupta, 2023). Variants of LDMR have been successfully applied to wind turbine fault detection (Tang et al., 2019), imbalanced classification (Cheng et al., 2016a,b), and structural health monitoring (Yang et al., 0).

However, when we examine the application of these kernel-based methods to wind power forecasting under extreme weather conditions, three critical gaps emerge:

Gap 1: Inability to handle distribution shift. Both SVR/SVQR and LDMR are predicated on the i.i.d. assumption—they assume that training and testing data are drawn from the same distribution. This assumption is fundamentally violated during extreme weather events, where the joint distribution $P(\mathbf{X}, Y)$ differs drastically between normal (source domain) and extreme (target domain) conditions. A model trained solely on abundant normal weather data will exhibit severe performance degradation when deployed under extreme conditions, regardless of how well the margin distribution is optimized. While transfer learning and domain adaptation techniques have been extensively studied in computer vision (Tao et al., 2012) and sensor drift compensation (Yan et al., 2022; Guo et al., 2022; Zhai et al., 2023; Se et al., 2023), their integration with kernel-based regression methods for wind power forecasting remains largely unexplored.

Gap 2: Conflict between symmetric variance regularization and asymmetric quantile estimation. The core principle of LDMR—minimizing the variance of prediction residuals—implicitly assumes a symmetric error distribution centered at zero. This is appropriate for conditional mean estimation but fundamentally incompatible with quantile regression. For example, when estimating the 95th percentile (upper prediction bound), the optimal residual distribution should be asymmetric, with 95% of errors being negative (under-predictions) and only 5% positive (over-predictions). Naively applying LDMR's symmetric variance minimization to quantile regression will force the error distribution toward symmetry, thereby degrading interval calibration and violating the target coverage probability. To the best of our knowledge, no existing work has addressed this theoretical tension between distributional robustness and quantile estimation within the LDMR framework.

Gap 3: Neglect of label information in transfer learning. Existing transfer learning approaches for wind power forecasting predominantly adopt unsupervised domain adaptation strategies, focusing exclusively on aligning feature

distributions (e.g., via Maximum Mean Discrepancy) while treating target domain labels—even when available in small quantities—as entirely unused supervision signals (Tao et al., 2012; Guo et al., 2022). This is particularly problematic for extreme weather forecasting, where scarce target samples represent high-value data points that capture critical domain-specific dynamics. A principled transfer learning framework should leverage both feature-space alignment and label-space supervision, strategically weighting target samples to guide model adaptation.

1.2. Our Contribution: TLQLDMR

To bridge these gaps, we propose the Transfer Learning Quantile Large Margin Distribution Machine for Regression (TLQLDMR), a principled extension of the classical LDMR paradigm specifically designed for cross-domain wind power forecasting under extreme weather conditions. Our approach preserves the theoretical strengths of kernel-based methods—convex optimization, generalization guarantees, and robustness on limited data—while introducing three key innovations:

Innovation 1: Asymmetric Variance Regularizer for Quantile-Aware Margin Distribution. We introduce a novel asymmetric variance regularization term that decomposes the residual variance into quantile-specific components. Unlike standard LDMR, which imposes symmetric variance minimization $\text{Var}(y_i - f(\mathbf{x}_i))$, our formulation applies differential regularization strengths to positive and negative residuals, thereby compressing overall error dispersion while preserving the requisite skewness for accurate quantile prediction. This reconciles the tension between distributional robustness (tight margin concentration) and quantile estimation (asymmetric error distribution), enabling TLQLDMR to achieve both superior point prediction accuracy and well-calibrated prediction intervals.

Innovation 2: Supervised Transfer Learning with Dual-Space Adaptation. We develop a unified transfer learning framework that integrates two complementary adaptation mechanisms. In the feature space, we employ Maximum Mean Discrepancy (MMD) minimization to align the source (normal weather) and target (extreme weather) distributions in the Reproducing Kernel Hilbert Space (RKHS), ensuring that the learned feature representation is domain-invariant (Tao et al., 2012). Crucially, in the label space, we introduce a Weighted Pinball Loss mechanism that assigns substantially higher weights to scarce target domain samples, treating them as high-value corrective anchors that guide the model toward extreme weather dynamics. This dual-space adaptation strategy enables TLQLDMR to leverage abundant source data for structural learning while prioritizing target data for domain-specific refinement, addressing the label neglect problem in existing unsupervised domain adaptation methods.

Innovation 3: Scalable Nyström-SGD Optimization for Large-Scale Deployment. The standard quadratic programming (QP) formulation of kernel-based transfer learning incurs $O(N^3)$ computational complexity due to the need to compute and invert the full kernel matrix, rendering it impractical for large-scale wind farm datasets with tens of thousands of samples. We propose a scalable optimization algorithm that combines Nyström approximation—which constructs an explicit low-rank feature mapping using $m \ll N$ landmark points—with stochastic gradient descent (SGD). This reduces complexity to $O(T \cdot N \cdot m)$, where T is the number of iterations, enabling training to complete within seconds on commodity hardware while maintaining prediction accuracy within 0.5% of the exact QP solution.

1.3. Experimental Validation and Key Findings

We conduct extensive experiments on a publicly available multi-site wind farm SCADA dataset (State Grid Renewable Energy, 2022), encompassing six wind farms with capacities ranging from 36 MW to 200 MW. Extreme weather events are rigorously defined using statistical thresholds (98th percentile wind speed, 2nd/98th percentile temperature) and physical constraints (cut-out wind speed > 25 m/s). Our evaluation compares TLQLDMR against three categories of baselines: (1) classical kernel-based methods (SVR, SVQR, LDMR), (2) ensemble methods (Random Forest, Gradient Boosting), and (3) state-of-the-art transfer learning approaches (TrAdaBoost, domain adaptation SVR). The results demonstrate:

- **Superior deterministic prediction within the kernel-based paradigm:** TLQLDMR achieves 12–18% RMSE reduction compared to non-transfer kernel baselines (SVR, SVQR, LDMR) and 6–9% improvement over transfer learning baselines on the wind farm exhibiting the most severe distribution shift. This validates that our innovations—asymmetric variance regularization and supervised transfer learning—effectively extend the LDMR framework to handle cross-domain scenarios.
- **State-of-the-art probabilistic interval prediction:** At the 95% confidence level, TLQLDMR is the only method that simultaneously achieves target-level coverage (94.8%) and the lowest Coverage Width-based Criterion (CWC), indicating optimal reliability-sharpness balance. Classical SVQR under-covers (88–91%)

due to distribution shift, while ensemble methods produce excessively wide intervals. This demonstrates the effectiveness of our asymmetric variance regularizer in maintaining quantile calibration under domain shift.

- **Robustness to measurement noise:** Under injected Gaussian noise with signal-to-noise ratios (SNR) ranging from 20 dB to 5 dB, TLQLDMR maintains stable interval coverage (93–95%) while baselines degrade to 85–88%. This robustness stems from the margin distribution optimization principle inherited from LDMR, which concentrates errors tightly even under noisy conditions.
- **Computational efficiency for practical deployment:** The Nyström-SGD algorithm reduces training time from 47 minutes (standard QP solver) to 8 seconds on a dataset with 50,000 samples, achieving a 350× speedup with negligible accuracy loss (<0.5% RMSE increase). This makes TLQLDMR practical for operational wind farm forecasting systems that require frequent model retraining.

Comprehensive ablation studies confirm the necessity of each component: removing transfer learning, asymmetric variance regularization, or MMD alignment degrades either deterministic accuracy, interval calibration, or both. Residual density analysis further reveals that TLQLDMR produces well-calibrated error distributions that closely match the target quantile, whereas baselines exhibit systematic bias.

1.4. Positioning and Scope

It is important to clarify the scope and positioning of our work. TLQLDMR is not designed to compete with large-scale deep learning models (e.g., LSTM, Transformer, CNN-based architectures) that have demonstrated impressive performance on massive wind power datasets with millions of samples (Wang et al., 2019; Mao et al., 2025; Ye et al., 2025a). Rather, our contribution lies in advancing the kernel-based regression paradigm—specifically the LDMR framework—to handle transfer learning scenarios that are prevalent in extreme weather forecasting. Kernel methods offer distinct advantages in settings where: (1) training data is limited (hundreds to tens of thousands of samples), (2) theoretical guarantees and interpretability are valued, (3) computational resources are constrained, and (4) convex optimization is preferred over non-convex deep learning. By demonstrating that classical kernel-based methods can be effectively extended with transfer learning capabilities, we provide practitioners with a principled alternative to deep learning that is particularly well-suited for extreme weather forecasting where data scarcity is the primary bottleneck.

The remainder of this paper is organized as follows. Section 2 reviews the theoretical foundations of SVQR, LDMR, and transfer learning with MMD, establishing the building blocks for our method. Section 3 presents the TLQLDMR formulation, including the asymmetric variance regularizer, supervised transfer framework, and scalable optimization algorithm. Section 4 describes the dataset, experimental setup, and comprehensive evaluation results across deterministic prediction, probabilistic interval forecasting, ablation studies, and computational efficiency analysis. Finally, Section ?? concludes the paper and discusses future research directions.

2. Related Work

This section briefly reviews the theoretical foundations related to the proposed method, including Support Vector Quantile Regression, Large Margin Distribution Machine for Regression, and Transfer Learning with Maximum Mean Discrepancy.

2.1. Support Vector Quantile Regression (SVQR)

Traditional Support Vector Regression (SVR) relies on the ϵ -insensitive loss function, which assumes a symmetric error distribution (e.g., Gaussian) and estimates the conditional mean of the target variable. However, in scenarios like partial shading or wind curtailment, the noise distribution is often asymmetric and heteroscedastic. To address this, Support Vector Quantile Regression (SVQR) was introduced to estimate conditional quantiles, providing a comprehensive view of the outcome distribution.

The core of SVQR is the non-symmetric pinball loss function, defined as:

$$L_{\tau}(u) = \begin{cases} \tau u, & u \geq 0 \\ (\tau - 1)u, & u < 0, \end{cases} \quad (1)$$

where $\tau \in (0, 1)$ indicates the target quantile. Unlike SVR, SVQR finds a decision boundary that divides the data such that a proportion τ of the samples lie below the boundary. The optimization problem minimizes the structural risk along with the empirical risk measured by the pinball loss:

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n L_{\tau}(y_i - (\mathbf{w}^T \phi(\mathbf{x}_i) + b)), \quad (2)$$

this formulation allows SVQR to capture heavy-tailed distributions and outliers more effectively than mean-based regression.

2.2. Large Margin Distribution Machine for Regression (LDMR)

While SVR and SVQR focus on maximizing the minimum margin (the distance from the closest instances to the hyperplane), they largely ignore the statistical distribution of the margin over the entire dataset. The **Large Margin Distribution Machine (LDMR)** was proposed to improve generalization by explicitly optimizing the distribution of the margins (or residuals).

The key idea of LDMR is to not only maximize the margin of the support vectors but also to minimize the variance of the error distribution for all samples. A typical LDMR formulation for regression integrates two additional statistical regularization terms into the objective function:

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 + \lambda_1 \text{Var}(\text{Residuals}) + \lambda_2 \text{Mean}(\text{Residuals}) + C \sum_{i=1}^n L(y_i, f(\mathbf{x}_i)), \quad (3)$$

by constraining the variance, LDMR forces the prediction errors to be clustered tightly, leading to a more compact and robust hypothesis space. Our work extends this concept by modifying the symmetric variance term of standard LDMR to an asymmetric one suitable for quantile regression.

2.3. Transfer Learning and Maximum Mean Discrepancy (MMD)

Standard machine learning models operate under the assumption that training (source) and testing (target) data are drawn from the same probability distribution. When this assumption is violated—such as in wind power forecasting across different weather regimes—the model performance degrades significantly. Transfer learning aims to transfer knowledge from a label-rich source domain to a target domain with different but related distributions.

A fundamental challenge in transfer learning is measuring and reducing the distribution discrepancy. Maximum Mean Discrepancy (MMD) is a widely used non-parametric metric that measures the distance between two distributions in a Reproducing Kernel Hilbert Space (RKHS). Given source data $\{\mathbf{x}_{S,i}\}$ and target data $\{\mathbf{x}_{T,j}\}$, the squared MMD distance is defined as:

$$\text{MMD}^2(D_S, D_T) = \left\| \frac{1}{n_S} \sum_{i=1}^{n_S} \phi(\mathbf{x}_{S,i}) - \frac{1}{n_T} \sum_{j=1}^{n_T} \phi(\mathbf{x}_{T,j}) \right\|_{\mathcal{H}}^2, \quad (4)$$

where $\phi(\cdot)$ is the feature mapping induced by the kernel. By minimizing the MMD distance during training, the model learns a feature representation where the source and target domains are statistically aligned, enabling the classifier or regressor trained on the source domain to generalize well to the target domain.

3. Proposed Methodology

In this section, we elaborate on the construction of the Transfer Learning Quantile Regression Model (TLQLDMR). We first define the necessary notations, then present the final objective function which integrates Structural Risk, Distribution Variance, Domain Adaptation (MMD), and Weighted Pinball Loss. Finally, we describe the constraints introduced to handle the non-differentiability of the Quantile Loss.

3.1. Notations

Consider the problem of wind power forecasting under extreme weather conditions. We define the **Source Domain (Normal Weather)** as $D_S = \{(\mathbf{x}_{S,i}, y_{S,i})\}_{i=1}^{n_S}$, containing abundant conventional weather samples, and the **Target Domain (Extreme Weather)** as $D_T = \{(\mathbf{x}_{T,j}, y_{T,j})\}_{j=1}^{n_T}$, containing scarce extreme weather samples. The input space

$\mathcal{X} \subseteq \mathbb{R}^d$ represents feature vectors constructed via sliding window processing, and the output space $\mathcal{Y} \subseteq \mathbb{R}$ represents wind power. The sample sizes satisfy $n_T \ll n_S$.

Due to drastic meteorological changes during extreme events, the joint distributions differ significantly: $P_S(\mathbf{X}_S, \mathbf{Y}_S) \neq P_T(\mathbf{X}_T, \mathbf{Y}_T)$. This distribution discrepancy implies that a model trained solely on P_S will fail on P_T . We define a feature mapping $\phi : \mathcal{X} \rightarrow \mathcal{H}$ that maps the input to a high-dimensional Reproducing Kernel Hilbert Space (RKHS). The prediction model is defined as:

$$f(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) + b, \quad (5)$$

where $\mathbf{w}$ is the weight vector and b is the bias. The quantile $\tau \in (0, 1)$ is set to a specific value (e.g., $\tau = 0.95$) for interval prediction (e.g., predicting the upper bound).

3.2. TLQLDMR Model Formulation

The overall architecture of the proposed TLQLDMR model is illustrated in Fig. 1. The model integrates four key components: structural risk minimization, asymmetric variance regularization, MMD-based domain adaptation, and weighted pinball loss for supervised transfer learning.

3.2.1. Objective Function

Our optimization goal is to find the optimal $\mathbf{w}, b$ and slack variables ξ, ξ^* to minimize the following objective function J . Let $\bar{u}_S = \frac{1}{n_S} \sum_{i=1}^{n_S} (y_{S,i} - f(\mathbf{x}_{S,i}))$ denote the mean residual over the source domain:

$$\begin{aligned} \min_{\mathbf{w}, b, \xi, \xi^*} \quad J = & \underbrace{\frac{1}{2} \|\mathbf{w}\|^2}_{(1) \text{ Structural Risk}} \\ & + \underbrace{\lambda_1 \sum_{i=1}^{n_S} ((y_{S,i} - \mathbf{w}^T \phi(\mathbf{x}_{S,i}) - b) - \bar{u}_S)^2}_{(2) \text{ Distribution Variance (LDMR Core)}} \\ & + \underbrace{\lambda_2 \left\| \frac{1}{n_S} \sum_{i=1}^{n_S} \phi(\mathbf{x}_{S,i}) - \frac{1}{n_T} \sum_{j=1}^{n_T} \phi(\mathbf{x}_{T,j}) \right\|_{\mathcal{H}}^2}_{(3) \text{ Domain Adaptation (MMD)}} \\ & + \underbrace{C_S \sum_{i=1}^{n_S} L_{\tau}(y_{S,i}, f(\mathbf{x}_{S,i})) + C_T \sum_{j=1}^{n_T} L_{\tau}(y_{T,j}, f(\mathbf{x}_{T,j}))}_{(4) \text{ Weighted Pinball Loss}}, \end{aligned} \quad (6)$$

the physical meaning of each term is as follows:

1. **Structural Risk:** This term controls the complexity of the model and prevents overfitting, embodying the maximum margin principle of SVM.
2. **Distribution Variance:** This term is the core of our proposed Quantile LDMR (QLDMR), incorporating a critical modification to the original LDMR variance.
 - **Theoretical Conflict:** The original LDMR minimizes the variance of absolute residuals, which assumes a symmetric error distribution suitable for mean regression. However, this symmetry contradicts the inherent asymmetry of Quantile Regression (governed by τ). Applying a symmetric variance penalty would force the residual distribution towards the center, degrading the skewness required for accurate quantile estimation.
 - **Asymmetric Modification:** To resolve this, we employ an **Asymmetric Variance** regularization based on the algebraic (signed) residuals. This design allows the model to reduce error dispersion (compactness) without constraining the distribution's direction, thereby preserving the non-symmetric properties essential for the Pinball Loss and improving interval coverage.




Fig. 1: The overall framework of the proposed TLQLDMR model. Source domain (normal weather) and target domain (extreme weather) data are mapped into a high-dimensional RKHS via kernel mapping. The joint optimization objective integrates structural risk, distribution variance, MMD-based domain adaptation, and weighted pinball loss. The model outputs probabilistic prediction intervals for robust wind power forecasting under extreme conditions.

- Domain Adaptation (MMD):** In the context of wind power forecasting under extreme weather conditions, our primary objective is to mitigate the differences in data distributions across various domains by means of learning a feature transformation. Consequently, the effective measurement of distribution disparities between domains becomes a paramount concern. In this pursuit, it is crucial to select a method suitable for quantifying distribution dissimilarity, while ensuring that the chosen method is computationally efficient and easily integrated

into the optimization problem of our model. Maximum Mean Discrepancy (MMD) has been demonstrated as an effective approach for estimating the distance between two distributions, offering a relatively straightforward implementation with fast computation speed. It targets the input space $\mathcal{X}$ to minimize the distance between the distribution centers of the source and target domains.

4. **Weighted Pinball Loss:** This term targets the output space $\mathcal{Y}$ and acts as the decisive component for “Supervised Transfer Learning”. Unlike unsupervised transfer which relies solely on feature alignment, we explicitly incorporate the target domain samples D_T into the loss calculation for the following reasons:

- **Corrective Anchors:** The target samples act as “anchors for correction”. The source domain data (D_S) teaches the model general laws (e.g., higher wind speed leads to higher power), but extreme weather samples (D_T) often violate these laws (e.g., high wind speed causing cut-out and zero power). If D_T were excluded from the loss, the model, driven implicitly by D_S , would wrongly predict full power during extreme events. Including D_T in the loss imposes a heavy penalty for such mispredictions, forcing the model to learn the specific mapping of extreme events.
- **Differentiation from MMD:** While MMD brings the “appearance” (features) of normal and extreme weather closer, the Pinball Loss ensures the “content” (labels) is accurate. It forces the regression curve to correctly envelope the rare data points that MMD cannot correct on its own.
- **Asymmetric Weighted Strategy:** We employ a weighted strategy to address the extreme sample imbalance ($n_T \ll n_S$). By setting $C_T \gg C_S$ (e.g., $C_T = 100C_S$), the optimization process assigns substantially higher importance to target domain samples. This weighting scheme ensures that prediction errors on scarce extreme weather data incur penalties commensurate with their practical significance, preventing them from being dominated by the abundant source domain samples during gradient updates.

3.2.2. Constraints

To handle the non-differentiability of the Pinball Loss, we introduce slack variables ξ_k, ξ_k^* to transform it into linear constraints. For every sample k in the training set $D = D_S \cup D_T$:

$$\begin{cases} y_k - (\mathbf{w}^T \phi(\mathbf{x}_k) + b) \leq \xi_k, & \forall k \in \{1, \dots, n_S + n_T\} \\ (\mathbf{w}^T \phi(\mathbf{x}_k) + b) - y_k \leq \xi_k^*, & \forall k \in \{1, \dots, n_S + n_T\} \\ \xi_k \geq 0, & \xi_k^* \geq 0, \end{cases} \quad (7)$$

in the objective function, the term $L_\tau(y_k, f(\mathbf{x}_k))$ corresponds to $\tau \xi_k + (1 - \tau) \xi_k^*$.

3.3. Model Optimization

In this section, we derive the dual problem of TLQLDMR to enable efficient solving using Quadratic Programming (QP).

3.3.1. Matrix Reformulation of the Primal Problem

To facilitate the derivation, we rewrite the variance term and the MMD term in the objective function [Eq. (6)] using matrix traces or norms. Let $X = [X_S, X_T]$ denote the entire training set with $N = n_S + n_T$ samples. For consistency with standard derivations, we add a coefficient of 1/2 to the variance and MMD terms. The primal objective can be reformulated as:

$$\min_{\mathbf{w}, b, \xi, \xi^*} \frac{1}{2} \|\mathbf{w}\|^2 + \frac{\lambda_1}{2} \text{Var}_S(\mathbf{w}, b) + \frac{\lambda_2}{2} \text{MMD}^2(\mathbf{w}) + \sum_{k=1}^N C_k L_\tau(y_k, f(\mathbf{x}_k)) \quad (8)$$

where C_k takes value C_S if sample $k \in D_S$ and C_T if $k \in D_T$.

We introduce two auxiliary matrices to express the variance and MMD terms in Eq. (8) as quadratic forms of the weights:

1. **Variance Matrix S:** To represent the source domain variance $\sum_{i \in D_S} (y_{S,i} - f(\mathbf{x}_{S,i}) - \bar{u}_S)^2$, we define a centering matrix $\mathbf{H}_S = \mathbf{I}_{n_S} - \frac{1}{n_S} \mathbf{1}_{n_S} \mathbf{1}_{n_S}^T$. The variance term can be written as $\|\mathbf{H}_S (\mathbf{y}_S - \mathbf{f}_S)\|^2$. To extend this to the full dataset, we define an $N \times N$ matrix $\mathbf{S}$:

$$\mathbf{S} = \begin{bmatrix} \mathbf{H}_S & \mathbf{0} \\ \mathbf{0} & \mathbf{0} \end{bmatrix} \quad (9)$$

2. **MMD Matrix $\mathbf{M}$:** The MMD term represents the distance between means. We define the MMD coefficient matrix $\mathbf{M} \in \mathbb{R}^{N \times N}$ as:

$$M_{ij} = \begin{cases} \frac{1}{n_S^2} & i, j \in D_S \\ \frac{1}{n_T^2} & i, j \in D_T \\ \frac{-1}{n_S n_T} & \text{otherwise} \end{cases} \quad (10)$$

3.3.2. Representer Theorem and Kernelization

According to the Representer Theorem [Schölkopf et al. \(2001\)](#), the optimal weight vector $\mathbf{w}$ lies in the span of the feature mappings of the training samples:

$$\mathbf{w} = \sum_{i=1}^N \beta_i \phi(\mathbf{x}_i) = \Phi \boldsymbol{\beta} \quad (11)$$

where $\Phi = [\phi(\mathbf{x}_1), \dots, \phi(\mathbf{x}_N)]$ is the feature matrix and $\boldsymbol{\beta} = [\beta_1, \dots, \beta_N]^T$ are the coefficients. Substituting Eq. (11) into the prediction function, we get $f(\mathbf{X}) = \boldsymbol{\beta}^T \mathbf{K} + b \mathbf{1}^T$, where $\mathbf{K} = \Phi^T \Phi$ is the $N \times N$ kernel matrix. Consequently, using Eq. (9) and Eq. (10), the structural risk, variance, and MMD terms can be expressed in terms of $\boldsymbol{\beta}$:

- Structural Risk: $\|\mathbf{w}\|^2 = \boldsymbol{\beta}^T \mathbf{K} \boldsymbol{\beta}$.
- Distribution Variance: The quadratic part with respect to $\boldsymbol{\beta}$ is $\boldsymbol{\beta}^T \mathbf{K} \mathbf{S} \mathbf{K} \boldsymbol{\beta}$. Note that the bias b is eliminated by the centering matrix $\mathbf{H}_S$.
- Domain Adaptation (MMD): $\text{MMD}^2(\mathbf{w}) = \boldsymbol{\beta}^T \mathbf{K} \mathbf{M} \mathbf{K} \boldsymbol{\beta}$.

Combining these, we define the total coefficient matrix $\boldsymbol{\Omega}$:

$$\boldsymbol{\Omega} = \mathbf{K} + \lambda_1 \mathbf{K} \mathbf{S} \mathbf{K} + \lambda_2 \mathbf{K} \mathbf{M} \mathbf{K} \quad (12)$$

By substituting the kernelized terms and Eq. (12) into Eq. (8), the kernelized primal problem becomes (ignoring constant terms w.r.t b in variance):

$$\min_{\boldsymbol{\beta}, b, \xi, \xi^*} \quad \frac{1}{2} \boldsymbol{\beta}^T \boldsymbol{\Omega} \boldsymbol{\beta} - (\lambda_1 \mathbf{K} \mathbf{S} \mathbf{Y})^T \boldsymbol{\beta} + \sum_{k=1}^N C_k (\tau \xi_k + (1 - \tau) \xi_k^*) \quad (13)$$

subject to the constraints from Eq. (7):

$$y_k - (\mathbf{K}_k \boldsymbol{\beta} + b) \leq \xi_k, \quad (\mathbf{K}_k \boldsymbol{\beta} + b) - y_k \leq \xi_k^*, \quad \xi_k, \xi_k^* \geq 0 \quad (14)$$

3.3.3. Lagrangian and Dual Problem

We verify the problem by introducing Lagrange multipliers $\alpha_k, \alpha_k^*, \mu_k, \mu_k^* \geq 0$ for Eq. (14). The Lagrangian function $\mathcal{L}$ is constructed as:

$$\begin{aligned} \mathcal{L} = & \frac{1}{2} \boldsymbol{\beta}^T \boldsymbol{\Omega} \boldsymbol{\beta} - (\lambda_1 \mathbf{K} \mathbf{S} \mathbf{Y})^T \boldsymbol{\beta} + \sum_{k=1}^N C_k (\tau \xi_k + (1 - \tau) \xi_k^*) \\ & - \sum_{k=1}^N \alpha_k (\xi_k - y_k + \mathbf{K}_k \boldsymbol{\beta} + b) - \sum_{k=1}^N \alpha_k^* (\xi_k^* + y_k - \mathbf{K}_k \boldsymbol{\beta} - b) - \sum (\mu_k \xi_k + \mu_k^* \xi_k^*) \end{aligned} \quad (15)$$

By setting the partial derivatives of Eq. (15) with respect to the primal variables $(b, \xi, \xi^*, \boldsymbol{\beta})$ to zero, we obtain the KKT conditions:

$$\frac{\partial \mathcal{L}}{\partial b} = 0 \implies \sum_{k=1}^N (\alpha_k - \alpha_k^*) = 0 \quad (16)$$

$$\frac{\partial \mathcal{L}}{\partial \xi_k} = 0, \frac{\partial \mathcal{L}}{\partial \xi_k^*} = 0 \implies \alpha_k \in [0, C_k \tau], \quad \alpha_k^* \in [0, C_k(1 - \tau)] \quad (17)$$

$$\frac{\partial \mathcal{L}}{\partial \beta} = 0 \implies \Omega \beta - \lambda_1 \mathbf{K} \mathbf{S} \mathbf{Y} - \mathbf{K}(\alpha - \alpha^*) = 0 \quad (18)$$

Let $\gamma = \alpha - \alpha^*$. From Eq. (18), the optimal β is derived as:

$$\beta = \Omega^{-1}(\mathbf{K} \gamma + \lambda_1 \mathbf{K} \mathbf{S} \mathbf{Y}) \quad (19)$$

Substituting Eq. (19) back into Eq. (15), and simplifying using Eq. (16) and Eq. (17), we arrive at the **Dual Problem**, which is a standard Quadratic Programming (QP) problem:

$$\begin{aligned} \min_{\alpha, \alpha^*} \quad & \frac{1}{2}(\alpha - \alpha^*)^T \mathbf{Q}_{dual}(\alpha - \alpha^*) + \mathbf{p}_{dual}^T(\alpha - \alpha^*) \\ \text{s.t.} \quad & \sum_{k=1}^N (\alpha_k - \alpha_k^*) = 0 \\ & 0 \leq \alpha_k \leq C_k \tau, \quad \forall k \\ & 0 \leq \alpha_k^* \leq C_k(1 - \tau), \quad \forall k \end{aligned} \quad (20)$$

where the dual matrix $\mathbf{Q}_{dual}$ and vector $\mathbf{p}_{dual}$ are defined as:

$$\mathbf{Q}_{dual} = \mathbf{K} \Omega^{-1} \mathbf{K}, \quad \mathbf{p}_{dual} = \lambda_1 \mathbf{K} \Omega^{-1} \mathbf{K} \mathbf{S} \mathbf{Y} - \mathbf{Y} \quad (21)$$

Once the optimal α, α^* are found using a QP solver for Eq. (20), we can recover β using Eq. (19) and construct the final decision function $f(\mathbf{x})$.

3.4. The Algorithm and Complexity Analysis

The complete training process of the proposed TLQLDMR model is summarized in Algorithm 1. The algorithm mainly consists of matrix construction and solving a Quadratic Programming (QP) problem.

The computational complexity of TLQLDMR primarily stems from three parts: matrix multiplication, matrix inversion, and the QP solver.

1. **Matrix Operations:** The construction of $\mathbf{K}$, $\mathbf{S}$, and $\mathbf{M}$ involves $O(N^2 d)$ or $O(N^2)$ operations. The calculation of Ω involves matrix multiplications of size $N \times N$, leading to a complexity of $O(N^3)$.
2. **Matrix Inversion:** Inverting the $N \times N$ matrix Ω requires $O(N^3)$ complexity.
3. **QP Solver:** Solving the standard Quadratic Programming problem for the dual variables typically has a complexity ranging from $O(N^{2.3})$ to $O(N^3)$ depending on the solver implementation (e.g., Interior Point Method).

Therefore, the overall time complexity of the proposed algorithm is approximately $O(N^3)$. Given that $N = n_S + n_T$ and $n_T \ll n_S$, the scale is dominated by the source domain sample size. While this complexity is acceptable for moderate-sized datasets, it becomes prohibitive for large-scale wind power forecasting scenarios involving tens of thousands of samples. To overcome this computational bottleneck and enable practical deployment in real-world wind farm applications, we develop a scalable optimization algorithm in the following subsection.

3.5. Scalable Optimization via Nyström Approximation

To address the computational bottleneck of the standard QP-based solver, we propose a scalable optimization algorithm that combines Nyström approximation with stochastic gradient descent. This approach reduces the time complexity from $O(N^3)$ to $O(T \cdot N \cdot m)$, where T denotes the number of iterations and $m \ll N$ is the number of landmark points.

Algorithm 1: TLQLDMR Training Algorithm

Input : Source Data $D_S = \{(\mathbf{x}_{S,i}, y_{S,i})\}_{i=1}^{n_S}$, Target Data $D_T = \{(\mathbf{x}_{T,j}, y_{T,j})\}_{j=1}^{n_T}$, Quantile τ , Kernel parameters, Regularization parameters λ_1, λ_2 , Pinball Loss weights C_S, C_T .

Output: Optimal dual variables α, α^* and bias b .

// Step 1: Data Preparation

- 1 Combine Source and Target data: $\mathbf{X} = [\mathbf{X}_S, \mathbf{X}_T]$, $\mathbf{Y} = [\mathbf{Y}_S; \mathbf{Y}_T]$, $N = n_S + n_T$;
- 2 Set weights vector $\mathbf{C} \in \mathbb{R}^N$ where $C_k = C_S$ for $k \leq n_S$ and C_T otherwise;

// Step 2: Matrix Construction

- 3 Compute Kernel Matrix $\mathbf{K} \in \mathbb{R}^{N \times N}$ using the specified kernel function;
- 4 Construct Variance Matrix $\mathbf{S}$ using Eq. (9);
- 5 Construct MMD Matrix $\mathbf{M}$ using Eq. (10);

// Step 3: Dual Problem Formulation

- 6 Compute coefficient matrix $\mathbf{\Omega} = \mathbf{K} + \lambda_1 \mathbf{K} \mathbf{S} \mathbf{K} + \lambda_2 \mathbf{K} \mathbf{M} \mathbf{K}$ using Eq. (12);
- 7 Invert $\mathbf{\Omega}$ to fit the dual form;
- 8 Compute Dual Matrix $\mathbf{Q}_{dual} = \mathbf{K} \mathbf{\Omega}^{-1} \mathbf{K}$ using Eq. (21);
- 9 Compute Dual Vector $\mathbf{p}_{dual} = \lambda_1 \mathbf{K} \mathbf{\Omega}^{-1} \mathbf{K} \mathbf{S} \mathbf{Y} - \mathbf{Y}$ using Eq. (21);

// Step 4: Solve QP

- 10 Solve the QP problem in Eq. (20) providing $\mathbf{Q}_{dual}$, $\mathbf{p}_{dual}$, and constraints $0 \leq \alpha_k \leq C_k \tau$, $0 \leq \alpha_k^* \leq C_k (1 - \tau)$;
- 11 **Return** Optimal α, α^* ;

// Step 5: Recover Primal Variables

- 12 Compute $\gamma = \alpha - \alpha^*$;
- 13 Recover β using Eq. (19);
- 14 Compute bias b using complementary slackness conditions (averaged over Support Vectors);

3.5.1. Nyström Feature Mapping

The core idea is to approximate the implicit infinite-dimensional feature mapping $\phi(\mathbf{x})$ with an explicit finite-dimensional representation. Given the training data, we first select m landmark points $Z = \{z_1, \dots, z_m\}$ from the source domain D_S using K-Means clustering to capture the underlying data manifold.

Let $\mathbf{W} \in \mathbb{R}^{m \times m}$ denote the kernel matrix of landmarks, where $W_{ij} = k(z_i, z_j)$. The explicit feature mapping $\psi : \mathcal{X} \rightarrow \mathbb{R}^m$ is then constructed as:

$$\psi(\mathbf{x}) = \mathbf{W}^{-1/2} [k(\mathbf{x}, z_1), \dots, k(\mathbf{x}, z_m)]^T \quad (22)$$

where $\mathbf{W}^{-1/2}$ is computed via eigendecomposition. This mapping satisfies the approximation property $k(\mathbf{x}_i, \mathbf{x}_j) \approx \psi(\mathbf{x}_i)^T \psi(\mathbf{x}_j)$, enabling us to work directly in the primal space.

3.5.2. Primal Formulation with Explicit Features

With the explicit feature mapping, the prediction function becomes $f(\mathbf{x}) = \mathbf{w}^T \psi(\mathbf{x}) + b$, where $\mathbf{w} \in \mathbb{R}^m$. Let $\mathbf{\Psi}_S \in \mathbb{R}^{n_S \times m}$ and $\mathbf{\Psi}_T \in \mathbb{R}^{n_T \times m}$ denote the mapped feature matrices for source and target domains, respectively. The regularization matrix $\mathbf{R} \in \mathbb{R}^{m \times m}$ incorporating structural risk, variance, and MMD terms is formulated as:

$$\mathbf{R} = \mathbf{I} + \underbrace{\lambda_1 \mathbf{\Psi}_S^T \mathbf{H}_S \mathbf{\Psi}_S}_{\text{Variance}} + \underbrace{\lambda_2 (\bar{\psi}_S - \bar{\psi}_T)(\bar{\psi}_S - \bar{\psi}_T)^T}_{\text{MMD}} \quad (23)$$

where $\mathbf{H}_S = \mathbf{I}_{n_S} - \frac{1}{n_S} \mathbf{1}_{n_S} \mathbf{1}_{n_S}^T$ is the centering matrix, and $\bar{\psi}_S, \bar{\psi}_T$ are the mean feature vectors of source and target domains.

3.5.3. Stochastic Gradient Descent Optimization

Instead of solving the dual QP problem, we directly optimize the primal objective using mini-batch stochastic gradient descent with the Adam optimizer. For each sample k in a mini-batch $\mathcal{B}$, the sub-gradient of the Pinball Loss

Algorithm 2: Fast TLQLDMR via Nyström Approximation

Input : Source data D_S , Target data D_T , Number of landmarks m , Quantile τ , Hyperparameters $\lambda_1, \lambda_2, C_S, C_T$, Learning rate η , Batch size B , Max iterations T .
Output: Optimal weights $\mathbf{w} \in \mathbb{R}^m$ and bias $b \in \mathbb{R}$.
// Step 1: Nyström Feature Mapping
1 Select m landmarks $Z = \{z_1, \dots, z_m\}$ from D_S via K-Means;
2 Compute landmark kernel matrix $\mathbf{W} \in \mathbb{R}^{m \times m}$;
3 Compute $\mathbf{W}^{-1/2}$ via eigendecomposition;
4 Construct explicit features $\psi(\mathbf{x})$ for all samples using Eq. (22);
// Step 2: Pre-compute Regularization Matrix
5 Compute feature matrices Ψ_S, Ψ_T and mean vectors $\bar{\Psi}_S, \bar{\Psi}_T$;
6 Construct regularization matrix $\mathbf{R}$ using Eq. (23);
// Step 3: Stochastic Optimization
7 Initialize $\mathbf{w} \leftarrow \mathbf{0}, b \leftarrow 0$;
8 **for** $t = 1$ **to** T **do**
9 Shuffle combined dataset $D = D_S \cup D_T$;
10 **for** each mini-batch $B \subset D$ **do**
11 Compute predictions: $\hat{y}_k = \mathbf{w}^T \psi(\mathbf{x}_k) + b, \forall k \in B$;
12 Compute sub-gradients g_k using Eq. (24);
13 Apply sample weights: $\tilde{g}_k \leftarrow C_k \cdot g_k$;
14 Compute gradients $\nabla_{\mathbf{w}} J, \nabla_b J$ using Eq. (25);
15 Update $\mathbf{w}, b$ using Adam optimizer with learning rate η ;
16 **end**
17 **end**
18 **return** $\mathbf{w}, b$;

with respect to the prediction is:

$$g_k = \frac{\partial L_\tau}{\partial \hat{y}_k} = \begin{cases} \tau - 1, & \text{if } y_k < \hat{y}_k \\ \tau, & \text{if } y_k \geq \hat{y}_k \end{cases} \quad (24)$$

The weighted sub-gradient is computed as $\tilde{g}_k = C_k \cdot g_k$, where $C_k \in \{C_S, C_T\}$ depends on the domain membership. The parameter gradients are then:

$$\nabla_{\mathbf{w}} J = \frac{1}{|B|} \sum_{k \in B} \tilde{g}_k \psi(\mathbf{x}_k) + \mathbf{R} \mathbf{w}, \quad \nabla_b J = \frac{1}{|B|} \sum_{k \in B} \tilde{g}_k \quad (25)$$

The complete training procedure is summarized in Algorithm 2.

3.5.4. Complexity Analysis of the Scalable Algorithm

The proposed scalable algorithm achieves significant computational savings:

1. **Landmark Selection:** K-Means clustering requires $O(N \cdot m \cdot I_{km})$ operations, where I_{km} is the number of K-Means iterations.
2. **Feature Mapping:** Computing $\mathbf{W}^{-1/2}$ costs $O(m^3)$, and mapping all N samples requires $O(N \cdot m \cdot d)$.
3. **Regularization Matrix:** Constructing $\mathbf{R}$ involves $O(n_S \cdot m^2)$ for the variance term and $O(m^2)$ for the MMD term.
4. **SGD Iterations:** Each iteration processes $O(N)$ samples with $O(m)$ operations per sample, yielding $O(T \cdot N \cdot m)$ total complexity.

The overall time complexity is $O(T \cdot N \cdot m + m^3)$. Since $m \ll N$ (typically $m = O(\sqrt{N})$), this represents a substantial reduction from the $O(N^3)$ complexity of the standard QP solver. Table 1 summarizes the comparison between the two algorithms.

Table 1
Complexity Comparison between Standard QP and Scalable Algorithm

Component	Standard QP	Scalable (Ours)
Kernel/Feature Computation	$O(N^2 d)$	$O(Nmd)$
Matrix Inversion	$O(N^3)$	$O(m^3)$
Optimization	$O(N^3)$	$O(TNm)$
Total	$O(N^3)$	$O(TNm + m^3)$

3.5.5. Theoretical Justification

The convergence of the proposed stochastic optimization algorithm is guaranteed under standard assumptions. Since the Pinball Loss is convex and Lipschitz continuous, and the regularization term $\frac{1}{2}\mathbf{w}^T \mathbf{R} \mathbf{w}$ is strongly convex when $\mathbf{R}$ is positive definite, the overall objective function is strongly convex. According to the convergence theory of SGD for strongly convex functions [Bottou et al. \(2018\)](#), the algorithm converges to the global optimum at a rate of $O(1/T)$.

Furthermore, the Nyström approximation error is bounded. Let $\tilde{k}(\mathbf{x}_i, \mathbf{x}_j) = \psi(\mathbf{x}_i)^T \psi(\mathbf{x}_j)$ denote the approximated kernel. Under mild conditions on the kernel eigenvalue decay, the approximation error satisfies:

$$\mathbb{E} [\|\mathbf{K} - \tilde{\mathbf{K}}\|_F] \leq O\left(\frac{N}{\sqrt{m}}\right) \quad (26)$$

This bound ensures that with a sufficient number of landmarks, the approximation quality is well-controlled, and the solution obtained by the scalable algorithm closely approximates that of the exact QP solver.

4. Experiments

In this section, we conduct a comprehensive experimental evaluation of the proposed TLQLDMR model for wind power forecasting under extreme weather conditions. The experiments are organized into three progressive stages: (1) extreme weather sample screening and domain shift quantification across multiple wind farms (Experiment 1); (2) deterministic point prediction comparison on the wind farm exhibiting the most severe distribution shift (Experiment 2); and (3) probabilistic interval prediction evaluation under a 95% confidence level (Experiment 3). We first describe the dataset and experimental setup in Sections 4.1 and 4.4, then present the results and analysis in Sections 4.5–4.7.

4.1. Dataset Description

The experiments are conducted on a publicly available wind farm SCADA dataset published by the State Grid Renewable Energy Generation Dataset ([State Grid Renewable Energy, 2022](#)). The dataset contains supervisory control and data acquisition (SCADA) records from six onshore wind farms with nominal capacities ranging from 36 MW to 200 MW. Each record is sampled at 15-minute intervals and includes the following meteorological and operational variables:

- Wind speed at multiple heights: hub height, 50 m, 30 m, and 10 m (m/s);
- Wind direction at hub height (degrees);
- Air temperature (°C);
- Atmospheric pressure (hPa);
- Relative humidity (%);
- Active power output (MW).

Table 2 summarizes the basic characteristics of the six wind farms. Each farm contains approximately 70,000 time-stamped records spanning multiple years of continuous operation.

Table 2

Summary of wind farm sites used in the experiments.

Farm ID	Site Name	Nominal Capacity (MW)	Total Samples
1	Wind farm site 1	99	~70,000
2	Wind farm site 2	200	~70,000
3	Wind farm site 3	99	~70,000
4	Wind farm site 4	66	~70,000
5	Wind farm site 5	36	~70,000
6	Wind farm site 6	96	~70,000

4.2. Feature Engineering and Data Preprocessing

The feature set comprises multi-height wind speeds, air temperature, atmospheric pressure, relative humidity, and cyclically encoded wind direction. Specifically, the raw wind direction angle θ is transformed into two orthogonal components $\sin(\theta)$ and $\cos(\theta)$ to eliminate the discontinuity at the $0^\circ/360^\circ$ boundary, which is a well-established practice in meteorological feature engineering (Wang et al., 2019).

A sliding window of length $W = 24$ (corresponding to a 6-hour lookback horizon at 15-minute resolution) is applied to construct supervised learning samples. For each window, the feature matrix $\mathbf{X}_t \in \mathbb{R}^{W \times d}$ is flattened into a vector $\mathbf{x}_t \in \mathbb{R}^{Wd}$, and the target variable y_t is the power output at the next time step. All features and the target variable are standardized to zero mean and unit variance using statistics computed from the training set. Missing values are imputed with the per-feature mean to ensure numerical stability for kernel-based models.

4.3. Extreme Weather Event Definition

A critical prerequisite for evaluating transfer learning under extreme conditions is the rigorous definition of what constitutes an “extreme weather event.” We adopt a rule-based identification scheme that combines statistical thresholds with physical domain knowledge. A sample is classified as belonging to the *target domain* (extreme weather) if it satisfies any of the following criteria:

1. **Statistical wind speed extreme:** Hub-height wind speed exceeds the 98th percentile of the farm’s historical distribution;
2. **Temperature extreme:** Air temperature falls below the 2nd percentile or exceeds the 98th percentile;
3. **Cut-out wind speed:** Hub-height wind speed exceeds 25 m/s, the typical cut-out threshold for modern wind turbines.

Notably, the power ramp criterion (rapid power fluctuations exceeding 5% of nominal capacity per 15-minute interval) is deliberately excluded from the extreme definition. This design choice avoids conflating control-induced power curtailment or dispatch adjustments with genuinely weather-driven anomalies. The union of the above three rules yields a target domain that is both physically interpretable and statistically sufficient for model training. All remaining samples constitute the *source domain* (normal weather conditions).

4.4. Experimental Setup

4.4.1. Evaluation Metrics

For deterministic point prediction (Experiment 2), we employ the following metrics. Let y_i and $\hat{y}_i$ denote the true and predicted power values, respectively, for $i = 1, \dots, n$:

- **Root Mean Squared Error (RMSE):** $\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$
- **Mean Absolute Error (MAE):** $\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$
- **Coefficient of Determination (R^2):** $R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$

For probabilistic interval prediction (Experiment 3), let $[L_i, U_i]$ denote the predicted interval at confidence level $1 - \alpha$:

- **Prediction Interval Coverage Probability (PICP):** $\text{PICP} = \frac{1}{n} \sum_{i=1}^n \mathbb{I}(L_i \leq y_i \leq U_i)$
- **Mean Prediction Interval Width (MPIW):** $\text{MPIW} = \frac{1}{n} \sum_{i=1}^n (U_i - L_i)$
- **Prediction Interval Normalized Average Width (PINAW):** $\text{PINAW} = \frac{\text{MPIW}}{\max(y) - \min(y)}$
- **Coverage Width-based Criterion (CWC):** A composite metric that penalizes PINAW when PICP falls below the nominal coverage $1 - \alpha$:

$$\text{CWC} = \text{PINAW} \times (1 + \gamma(\text{PICP}) \cdot e^{-\eta(\text{PICP} - (1 - \alpha))}) \quad (27)$$

where $\gamma(\text{PICP}) = 1$ if $\text{PICP} < 1 - \alpha$ and 0 otherwise, and η is a penalty coefficient (set to 50 in our experiments);

- **Winkler Score:** An interval scoring rule that combines interval width with a penalty proportional to the magnitude of coverage violations:

$$\text{Winkler} = \frac{1}{n} \sum_{i=1}^n \begin{cases} (U_i - L_i) + \frac{2}{\alpha}(L_i - y_i), & y_i < L_i \\ (U_i - L_i), & L_i \leq y_i \leq U_i \\ (U_i - L_i) + \frac{2}{\alpha}(y_i - U_i), & y_i > U_i \end{cases} \quad (28)$$

PICP measures reliability (whether the interval achieves the target coverage), while MPIW and PINAW measure sharpness (how tight the interval is). An ideal model achieves $\text{PICP} \geq 1 - \alpha$ with the smallest possible PINAW.

4.4.2. Domain Shift Metrics

To quantify the distribution discrepancy between normal and extreme domains (Experiment 1), we employ two complementary metrics computed on both wind speed and power dimensions:

- **Wasserstein-1 distance (W_1):** The earth mover's distance between the marginal distributions of the source and target domains, measuring the minimum "transportation cost" to transform one distribution into the other;
- **Kolmogorov-Smirnov statistic (KS):** The supremum of the absolute difference between the empirical cumulative distribution functions of the two domains.

4.4.3. Implementation Details

All experiments are implemented in Python using NumPy and scikit-learn. For SVM-based models, the training sample size is capped at 4,000 source domain samples and 1,200 target domain samples to maintain tractable computation times for kernel methods. Tree-based and ensemble models use the full training set without subsampling. The TLQLDMR model employs a Nyström approximation with $m = 1,200$ landmark points and is optimized via stochastic gradient descent with a learning rate of 0.01 for 64 epochs and a mini-batch size of 256. All random splits use a fixed seed for reproducibility.

4.5. Experiment 1: Extreme Sample Screening and Domain Shift Quantification

The first experiment identifies extreme-weather samples across all six wind farms using the criteria in Section 4.3, and then quantifies source-target distribution shift for each farm. Table 3 reports the domain-shift statistics sorted by the composite Wasserstein distance $W_1(\text{wind}) + W_1(\text{power})$. Preliminary sensitivity checks show that a stricter 99th-percentile wind-speed rule leaves too few target-domain samples for stable model development, whereas the adopted 98th-percentile rule preserves event extremity while maintaining statistical sufficiency.

Fig. 2 visualizes the distribution shift between normal and extreme weather domains for all six wind farms. Each subplot displays the kernel density estimates of wind speed (upper panel) and power output (lower panel) for both domains. The visual comparison confirms the quantitative findings in Table 3: Farm 6 exhibits the most pronounced separation between the two distributions, particularly in the power dimension, where the extreme domain shows a bimodal pattern with substantial mass near zero power (cut-out events) and high power (pre-cut-out high wind conditions).

Several observations emerge from Table 3 and Fig. 2. First, all six farms exhibit non-trivial distribution shifts between normal and extreme domains, with KS statistics ranging from 0.116 to 0.460, confirming that extreme

Table 3

Domain shift metrics between normal and extreme weather domains across six wind farms. Farms are sorted by the composite Wasserstein distance in descending order.

Wind Farm	Extreme Samples	$W_1(\text{Wind})$	$W_1(\text{Power})$	KS(Wind)	KS(Power)
Farm 6 (96 MW)	4,144	4.060	21.057	0.389	0.344
Farm 3 (99 MW)	4,144	3.389	18.915	0.338	0.265
Farm 2 (200 MW)	4,211	3.467	16.166	0.333	0.142
Farm 5 (36 MW)	3,039	4.195	9.851	0.460	0.363
Farm 4 (66 MW)	4,046	4.235	7.759	0.347	0.201
Farm 1 (99 MW)	4,184	3.140	5.964	0.335	0.116

weather events induce statistically significant distributional changes in both wind speed and power output. Second, the magnitude of the power distribution shift varies substantially across farms: Farm 6 exhibits a Wasserstein distance of 21.057 in the power dimension, approximately 3.5 times that of Farm 1 (5.964). This disparity reflects differences in turbine control strategies, terrain effects, and the local meteorological characteristics of each site.

Farm 6 (Nominal capacity 96 MW) ranks first in the composite domain shift metric ($W_1(\text{wind}) + W_1(\text{power}) = 25.117$), indicating the most pronounced distributional discrepancy between its normal and extreme operating regimes. Consequently, **Farm 6 is selected as the target wind farm for Experiments 2 and 3**, providing the most stringent test bed for evaluating the robustness of transfer learning methods under extreme weather conditions.




Fig. 2: Distribution shift visualization for all six wind farms. From left to right, top to bottom: Site 1 (99 MW), Site 2 (200 MW), Site 3 (99 MW), Site 4 (66 MW), Site 5 (36 MW), and Site 6 (96 MW). Each subplot shows kernel density estimates of wind speed (upper panel) and power output (lower panel) for normal (source) and extreme (target) domains. Site 6 exhibits the most severe domain shift in the power dimension.

4.6. Experiment 2: Deterministic Point Prediction under Extreme Weather

The second experiment evaluates deterministic point prediction on Farm 6 and compares TLQLDMR with six baselines under the same data split and preprocessing protocol.

The data partition for Farm 6 is as follows:

- **Source domain** (normal weather): 66,009 samples;
- **Target domain** (extreme weather): 4,143 samples, further split into:
 - Training set: 3,728 samples (90%);
 - Test set: 415 samples (10%).

The target domain is split via random shuffling with a fixed seed to ensure reproducibility. The 90%/10% target split is selected to maximize training signal under scarce extreme samples while retaining an independent hold-out subset for evaluation. All baseline models are trained exclusively on target-domain training data, whereas TLQLDMR additionally leverages source-domain samples through transfer learning.

We compare TLQLDMR against six state-of-the-art non-deep-learning models published in 2025, spanning SVM variants, tree-based ensembles, and random forest methods specifically designed for wind power forecasting:

1. **HHO-SVR** (Houssein et al., 2025): Support Vector Regression with hyperparameters optimized by the Harris Hawks Optimization metaheuristic;
2. **FLSVR** (Meena et al., 2025): A functional iterative method for solving the Lagrangian dual of SVR, offering improved convergence properties;
3. **ARA-SVR** (Duan et al., 2025): SVR with Associated Road Analysis for feature selection, originally proposed for short-term traffic flow prediction;
4. **KMeans-GBT** (Rizkallah, 2025): A two-stage model combining KMeans clustering with Gradient Boosted Trees to capture heterogeneous input–output relationships;
5. **RF-WPF** (Mustaffa & Sulaiman, 2025): A Random Forest model tailored for wind power forecasting with optimized tree depth and feature subsampling;
6. **BRF-WPF** (Chen & Shi, 2025): Broad Random Forest integrating Broad Learning System feature expansion with Random Forest regression for short-term wind power prediction.

All baseline models use hyperparameters consistent with their respective publications or default configurations, without extensive grid search, to ensure a fair comparison that reflects each method’s out-of-the-box performance.

The key hyperparameters of TLQLDMR are determined via a structured search on the validation set (20% held out from the target training data). The optimal configuration is: $\lambda_1 = 5 \times 10^{-4}$, $\lambda_2 = 1 \times 10^{-4}$, $C_S = 0.05$, $C_T = 150$, $\gamma_{\text{kernel}} = 0.002$, Nyström components $m = 1,200$, learning rate $\eta = 0.01$, epochs = 64, batch size = 256, and $\tau = 0.5$ (median regression for point prediction). The large ratio $C_T/C_S = 3,000$ reflects the asymmetric weighting strategy that assigns substantially higher importance to the scarce extreme weather samples.

Table 4 presents the test set performance of all models on the extreme weather domain of Farm 6.

Table 4 and Figs. 3–4 jointly demonstrate that TLQLDMR provides the most accurate and stable point forecasts on the Farm 6 extreme-weather test set. Several observations are noteworthy:

TLQLDMR improves on the strongest non-transfer baseline (HHO-SVR), raising R^2 from 0.9537 to 0.9553 and reducing RMSE from 7.603 to 7.476. Although the numerical margin is moderate, it is obtained under the most severe domain-shift setting with only 4,143 extreme-weather samples. The curve-level evidence in Figs. 3–4 is consistent with the aggregate metrics: TLQLDMR tracks abrupt ramps and low-power transitions more closely, while several baselines show larger local deviations around rapid fluctuations. Among baselines, HHO-SVR remains the most competitive kernel method, RF-WPF remains the second-best overall alternative, and BRF-WPF does not surpass RF-WPF under the present small-sample extreme-weather regime. The observed gain of TLQLDMR is consistent with the combined effect of source-domain transfer, asymmetric target weighting, and cyclic wind-direction encoding. Because many extreme-weather periods contain near-zero power outputs, MAPE is numerically unstable and therefore RMSE and R^2 are treated as the principal deterministic indicators.

To assess the robustness of the models against measurement noise, which is prevalent in SCADA systems during extreme weather, we evaluate the point prediction performance by injecting Gaussian noise into both the training and

Table 4

Deterministic point prediction results on the extreme weather test set of Farm 6 (96 MW). Best results are highlighted in bold.

Model	RMSE	MAE	R^2
HHO-SVR	7.603	5.530	0.9537
FLSVR	8.794	6.301	0.9381
ARA-SVR	9.246	5.954	0.9316
KMeans-GBT	9.620	6.695	0.9259
RF-WPF	7.967	5.344	0.9492
BRF-WPF	9.719	6.870	0.9244
TLQLDMR	7.476	5.348	0.9553

Table 5

Model performance (R^2) under varying Signal-to-Noise Ratios (SNR) on the extreme weather test set. Best results are highlighted in bold.


Model	SNR = 10 dB	SNR = 20 dB	SNR = 30 dB	SNR = 40 dB	SNR = 60 dB
HHO-SVR	0.9309	0.9539	0.9542	0.9538	0.9537
FLSVR	0.9253	0.9371	0.9380	0.9381	0.9381
ARA-SVR	0.9145	0.9303	0.9311	0.9316	0.9316
KMeans-GBT	0.9137	0.9247	0.9209	0.9230	0.9262
RF-WPF	0.9265	0.9438	0.9471	0.9485	0.9469
BRF-WPF	0.8939	0.9144	0.9177	0.9250	0.9253
TLQLDMR	0.9419	0.9547	0.9584	0.9560	0.9603

testing input features. We simulate varying Signal-to-Noise Ratios ($\text{SNR} \in \{60, 40, 30, 20, 10\}$ dB). For this analysis, TLQLDMR utilizes a higher-dimensional Nyström mapping ($m = 1,500$) and longer training duration (80 epochs) to accommodate the increased input variance.


The R^2 results under different noise levels are summarized in Table 5.

The observation confirms that while all models experience performance degradation as the noise level increases (SNR decreases), TLQLDMR consistently maintains the highest R^2 across all testing conditions. In particular, TLQLDMR remains above 0.94 even at 10 dB, while BRF-WPF exhibits the steepest degradation, indicating higher sensitivity to measurement perturbations. The transfer learning framework, reinforced by Large Margin Distribution regularization, effectively mitigates overfitting to noisy target samples and demonstrates the strongest noise resilience among the comparative models.


(a) HHO-SVR

HHO-SVR vs True (Wind farm site 6 (Nominal capacity-96MW))
Segment [0, 415]

(b) FLSVR

FLSVR vs True (Wind farm site 6 (Nominal capacity-96MW))
Segment [0, 415]

(c) ARA-SVR

ARA-SVR vs True (Wind farm site 6 (Nominal capacity-96MW))
Segment [0, 415]

(d) KMeans-GBT








KMeans-GBT vs True (Wind farm site 6 (Nominal capacity-96MW))
Segment [0, 415]

Fig. 3: Prediction-curve comparison on the Farm 6 extreme-weather test set (Part I). One row corresponds to one model, and each row contains the full segment plus peak/valley zoom views.

(e) RF-WPF

RF-WPF vs True (Wind farm site 6 (Nominal capacity-96MW))
Segment [0, 415)

(f) BRF-WPF

BRF-WPF vs True (Wind farm site 6 (Nominal capacity-96MW))
Segment [0, 415)

(g) TLQLDMR




TL-QLDMR vs True (Wind farm site 6 (Nominal capacity-96MW))
Segment [0, 415)

Fig. 4: Prediction-curve comparison on the Farm 6 extreme-weather test set (Part II). One row corresponds to one model, and each row contains the full segment plus peak/valley zoom views.

4.7. Experiment 3: Probabilistic Interval Prediction under Extreme Weather

The third experiment extends evaluation from point prediction to probabilistic interval forecasting. The objective is to verify whether TLQLDMR can achieve reliable 95% prediction intervals while keeping interval width sufficiently sharp under extreme-weather volatility.

To support the calibration-based interval prediction framework, the target domain data is partitioned into four subsets:

- Training set: 2,983 samples (for model fitting);
- Validation set: 559 samples (for hyperparameter selection and width scaling);
- Calibration set: 745 samples (for Conformalized Quantile Regression adjustment);
- Test set: 415 samples (for final evaluation).

All models employ a two-stage calibration procedure based on Conformalized Quantile Regression (CQR):

1. **Conformal adjustment:** The model first predicts intervals on the calibration set. The conformal quantile $\hat{q}$ is computed as the $(1 - \alpha)$ -quantile of the nonconformity scores $\max(L_i - y_i, y_i - U_i)$, ensuring finite-sample coverage guarantees;
2. **Width scaling:** A scaling factor is selected on the validation set to optimize the trade-off between coverage and interval width, subject to the constraint that PICP approaches the target level $1 - \alpha + \delta$, where $\delta = 0.02$ is a safety margin.

All predicted intervals are clipped to the physical range $[0, P_{\text{rated}}]$, where P_{rated} is the nominal capacity of the wind farm.

The interval prediction baselines are organized into two categories:

SVM-based quantile regression variants:

1. **TSVQR** (Ye et al., 2023): Twin Support Vector Quantile Regression, which constructs separate hyperplanes for upper and lower quantiles;
2. **UQSVM-SSVQR** (Wang & Xu, 2024): Sparse Support Vector Quantile Regression with approximate kernel methods for scalability;
3. **NFS-SVQR** (Ye et al., 2025b): Nonlinear Feature Selection combined with SVQR, published in *Neural Networks* (2025);
4. **NuSVR-CI**: ν -Support Vector Regression with CQR-calibrated prediction intervals as a classical SVM baseline.

Recent wind power interval prediction models:

1. **HybridDL-Interval** (Mao et al., 2025): A hybrid deep learning framework for day-ahead wind power interval forecasting, combining temporal attention and quantile loss (approximate implementation);
2. **AMQRNN** (Ye et al., 2025a): A recurrent attention encoder-decoder network for multi-step interval wind power prediction, implemented using an attention-GRU network;
3. **QGBR** (Hu & Liu, 2025): Quantile Gradient Boosted Regression, representing the improved quantile ensemble learning framework for wind power prediction;
4. **CQR-MLP** (Hu et al., 2025): MLP quantile ensemble with dynamic conformalized quantile regression, acting as a simplified baseline for neural ensemble search methods.

For interval prediction, TLQLDMR trains two separate quantile models at τ_{low} and τ_{high} to construct the lower and upper bounds of the prediction interval. The quantile pair ($\tau_{\text{low}} = 0.2, \tau_{\text{high}} = 0.8$) is jointly selected with the hyperparameters on the validation set, using a coverage-priority criterion that first ensures PICP approaches the target coverage $0.95 + \delta$ and then minimizes PINAW via a learned interval scaling factor (width_scale). The resulting width scaling for TLQLDMR is 0.88, which yields reliable coverage while avoiding unnecessary interval inflation.

Table 6 presents the interval prediction results on the extreme weather test set of Farm 6 at the 95% confidence level ($\alpha = 0.05$).

The interval prediction results indicate a clear reliability-sharpness trade-off under extreme weather, and TLQLDMR offers the most favorable balance. TLQLDMR reaches the nominal 95% coverage target (PICP = 0.959)

Table 6

Interval prediction results on the extreme weather test set of Farm 6 at 95% confidence level ($\alpha = 0.05$). Baseline models are sorted by CWC in ascending order, while TLQLDMR is placed in the last row for direct comparison. Best results per column are highlighted in bold.

Model	PICP	MPIW	PINAW	CWC	Winkler
NuSVR-CI	0.843	19.970	0.208	0.285	54.603
HybridDL-Interval	0.954	31.645	0.329	0.329	60.224
TSVQR	0.889	24.592	0.256	0.342	48.178
QGBR	0.949	31.715	0.330	0.429	37.643
CQR-MLP	0.942	31.750	0.330	0.431	46.772
AMQRNN	0.940	32.432	0.337	0.441	66.874
NFS-SVQR	0.913	68.359	0.711	0.941	76.997
UQSVM-SSVQR	0.928	81.710	0.850	1.117	85.162
TLQLDMR	0.959	25.791	0.268	0.268	34.506

while obtaining the best composite quality (CWC = 0.268) and the best Winkler score (34.506), which shows that transfer learning with validation-based width scaling can provide both reliable and sharp intervals. By contrast, NuSVR-CI and TSVQR produce narrow intervals but severe under-coverage, whereas HybridDL-Interval, QGBR, CQR-MLP, and AMQRNN can approach acceptable coverage only with wider intervals. This comparison suggests that explicit asymmetric regularization combined with cross-domain adaptation is more effective for heavily skewed extreme-weather samples.

To assess interval robustness under sensor perturbation, Gaussian noise is injected only into test inputs at SNR $\in \{60, 40, 30\}$ dB, while calibration parameters remain fixed to preserve comparability across models. For TLQLDMR, a conservative scaling multiplier (1.15) is used to maintain coverage. Unlike the previous single-metric presentation, Table 7 reports the full set of interval metrics.

Across all noise levels, TLQLDMR maintains the best CWC while keeping PICP above 0.95. The CWC values remain tightly concentrated around 0.278, indicating stable interval quality under moderate measurement perturbation. In contrast, NuSVR-CI yields narrow intervals (small MPIW/PINAW) but exhibits severe under-coverage, and several deep/ensemble baselines require substantially wider intervals to approach reliable coverage. These results indicate that TLQLDMR preserves a stronger reliability-sharpness balance under measurement perturbations.

Table 7

Full interval metrics under injected test-input noise on the Farm 6 extreme-weather test set. For PICP, values near 0.95 are preferred; for MPIW, PINAW, CWC, and Winkler, lower values are preferred.

Model	SNR (dB)	PICP	MPIW	PINAW	CWC	Winkler
TSVQR	30	0.8892	24.570	0.2556	0.3422	48.306
	40	0.8892	24.588	0.2558	0.3424	48.140
	60	0.8892	24.592	0.2558	0.3425	48.169
UQSVM-SSVQR	30	0.9229	81.706	0.8499	1.1191	85.364
	40	0.9277	81.718	0.8501	1.1167	85.178
	60	0.9277	81.710	0.8500	1.1166	85.162
NFS-SVQR	30	0.9060	68.369	0.7112	0.9442	77.305
	40	0.9133	68.376	0.7113	0.9409	76.878
	60	0.9133	68.359	0.7111	0.9407	76.983
NuSVR-CI	30	0.8386	19.970	0.2077	0.2856	54.922
	40	0.8361	19.970	0.2077	0.2860	54.623
	60	0.8434	19.970	0.2077	0.2849	54.588
QGBR	30	0.9518	32.067	0.3336	0.3336	38.441
	40	0.9518	31.721	0.3300	0.3300	37.424
	60	0.9470	31.706	0.3298	0.4294	37.750
CQR-MLP	30	0.9325	32.548	0.3386	0.4438	48.828
	40	0.9325	32.550	0.3386	0.4438	48.780
	60	0.9325	32.554	0.3386	0.4438	48.743
HybridDL-Interval	30	0.9422	32.304	0.3360	0.4384	53.462
	40	0.9422	32.307	0.3361	0.4385	53.420
	60	0.9422	32.306	0.3361	0.4385	53.408
AMQRNN	30	0.9518	33.517	0.3487	0.3487	68.546
	40	0.9518	33.513	0.3486	0.3486	68.517
	60	0.9518	33.512	0.3486	0.3486	68.515
TLQLDMR	30	0.9542	26.709	0.2778	0.2778	36.307
	40	0.9566	26.713	0.2779	0.2779	36.034
	60	0.9566	26.715	0.2779	0.2779	36.083

4.8. Experiment 4: Ablation Study

To validate the contribution of the three core components (transfer learning, LDMR variance regularization, and MMD feature alignment), this ablation experiment evaluates structurally reduced variants on both point and interval tasks.

The full TLQLDMR framework is compared against three architectural reductions:

1. **w/o Transfer (Target-only QLDMR)**: Trained exclusively on extreme weather target samples without source domain integration or MMD alignment.
2. **w/o LDMR (Variance Penalty)**: The distribution variance regularizer is removed, deteriorating the framework into a generalized Transfer SVR/SVQR variant.
3. **w/o MMD (Feature Alignment)**: The algorithm employs cross-domain training data but disables Maximum Mean Discrepancy alignment, relying purely on weighted objective loss.

Table 8 reports the ablation results on Farm 6. The full TLQLDMR consistently attains the best deterministic and probabilistic performance, indicating that all three components contribute non-trivially to final accuracy and calibration quality.

The ablation outcomes confirm that each component contributes materially to robustness. Removing transfer learning reduces both point accuracy (R^2 : 0.9553 $\rightarrow$ 0.9497) and interval quality (CWC: 0.2683 $\rightarrow$ 0.2797), indicating that target-only training is insufficient under strong domain shift. Removing LDMR regularization or MMD alignment causes clear interval-quality degradation (CWC rises to 0.3507 and 0.3528), even when point-metric changes are

Table 8

Ablation study evaluations for point and interval extreme-weather prediction. Best results are highlighted in bold.

Model Variants	Point Prediction		Interval Prediction (95% Level)			
	RMSE	R^2	PICP	MPIW	PINAW	CWC
w/o Transfer	7.930	0.9497	0.952	26.891	0.2797	0.2797
w/o LDMR	8.205	0.9461	0.949	25.925	0.2697	0.3507
w/o MMD	7.745	0.9520	0.947	26.049	0.2710	0.3528
TLQLDMR (Full)	7.476	0.9553	0.959	25.791	0.2683	0.2683




Fig. 5: Residual density comparison between TLQLDMR and ablated variants on the Farm 6 test set.

smaller, showing that residual-shape control and feature alignment are central to reliable uncertainty estimation. The residual-density comparison in Fig. 5 supports this conclusion by showing tighter central concentration and weaker heavy tails for the full TLQLDMR model; some ablated variants exhibit sharper central peaks but noticeably thicker tails, which corresponds to a higher probability of large-error events.

4.9. Experiment 5: Scalable Optimization Validation

Experiment 5 evaluates whether the proposed Nyström + SGD solver can preserve practical training efficiency as data scale increases, compared with three alternative optimization strategies implemented in Train_TL_QLDMR: (i) full dual QP via CVXOPT, (ii) gradient-based QP via PyTorch (torch_gd), and (iii) kernel-slice mini-batch SGD (batch_gsd).

The combined training size N is progressively increased from 100 to 2,000. For each configuration, wall-clock training time (seconds) is recorded. Table 9 summarizes the results.

Both CVXOPT and torch_gd become infeasible beyond $N = 800$ due to cubically growing memory and computation demands. The kernel-slice batch_gsd avoids full matrix inversion but still requires $O(N^2)$ kernel evaluations per batch, yielding runtimes of 3–6 seconds. In contrast, the proposed Nyström + SGD solver maintains near-linear scaling and completes training for $N = 2,000$ in only 1.67 seconds, confirming the practical viability of the approximate primal formulation for large-scale wind farm datasets.

Table 9

Training time (seconds) of four solvers over expanding dataset scales. “—” denotes infeasible runs due to excessive resource consumption.

N	CVXOPT	torch_gd	batch_sgd	Nyström + SGD (Ours)
100	0.494	2.254	2.972	0.024
200	0.662	2.398	2.300	0.554
400	0.884	2.508	4.374	0.696
600	2.728	2.668	4.874	0.824
800	4.949	2.375	3.890	0.924
1,200	—	—	6.049	1.114
1,600	—	—	6.123	1.465
2,000	—	—	—	1.670

4.10. Discussion

The extended experimental evaluations yield the following conclusive insights:

Transfer learning improves deterministic forecasting under severe shift. Experiment 2 shows consistent gains in RMSE and R^2 , and curve-level evidence confirms better tracking during abrupt peak-valley transitions.

Interval forecasting achieves a better reliability-sharpness balance. In Experiment 3, TLQLDMR is the only method that simultaneously maintains target-level coverage and the best composite penalty (CWC), and this advantage persists under injected measurement noise.

Ablation confirms the necessity of each component. Removing transfer, LDMR variance regularization, or MMD alignment degrades either deterministic accuracy, interval calibration, or both; residual-density plots provide complementary evidence for this degradation.

The scalable optimizer materially reduces computational burden. Experiment 5 verifies that the Nyström + SGD formulation remains practical in regimes where full QP optimization fails due to computational constraints.

Limitations and future work. The present evaluation is based on onshore wind farms and fixed feature engineering pipelines; extension to offshore assets, dynamic re-calibration strategies, and hybrid deep-kernel transfer formulations remains a relevant direction.

5. Conclusion

Accurate wind power prediction under extreme meteorological conditions is crucial yet methodologically arduous due to extreme sample imbalance, drastic distribution shifts, and the probabilistic necessity of capturing highly asymmetric tail volatility. To overcome the inherent deficiencies of symmetric large margin algorithms and conventional unsupervised feature alignment techniques, this paper proposed the Transfer Learning Quantile Large Margin Distribution Machine for Regression (TLQLDMR).

Extensive evaluations on public operational wind farm SCADA datasets rigorously validated the TLQLDMR framework’s deterministic accuracy and probabilistic interval forecasting competence. First, its dual-space domain adaptation framework—unifying MMD-based feature alignment with an asymmetric weighted Pinball Objective—demonstrated unambiguous superiority over contemporaneous baseline algorithms, yielding definitive point prediction dominance under volatile measurement noise settings. Second, through carefully tuned interval scaling validation, TLQLDMR systematically achieved nominal 95% prediction interval coverage constraints while preserving the sharpest aggregate calibration boundaries (PINAW = 0.268, CWC = 0.268) amidst severe meteorological conditions. Third, systematic ablation analyses isolated and confirmed that the holistic fusion of cross-domain knowledge transposition, distributional variance compacting, and feature proximity alignment is strictly indispensable for managing heavy-tailed prediction fluctuations. Lastly, implementing a Nyström-approximated Stochastic Gradient Descent optimization routine conclusively dissolved typical kernel complexity limitations, fostering manageable scalability on vast sensor data histories previously intractable for rigorous analytical quadratic programming.

Looking forward, the extension of the proposed asymmetric variance constraints to broader spatiotemporal correlation tasks and its integration with sophisticated deep self-attention encoders present promising avenues toward fully adaptive sustainable energy intelligence.

CRedit authorship contribution statement

***: Methodology, Software, Writing - original draft. ***: Visualization, Data curation. ***: Formal analysis, Software. ***: Writing – review & editing. ***: Funding acquisition, Validation, Supervision.

Declaration of Competing Interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Acknowledgments

This work is supported by the National Key Research and Development Program of China (No. 2021YFB3101500) and the Chongqing Municipal Training Program of Innovation and Entrepreneurship for Undergraduate (No. X202310635204).

References

- Bottou, L., Curtis, F. E., & Nocedal, J. (2018). Optimization methods for large-scale machine learning. *SIAM Review*, 60, 223–311. doi:10.1137/16M1080173.
- Chen, Y., & Shi, J. (2025). Broad random forest: A lightweight prediction model for short-term wind power by fusing broad learning and random forest. *Sustainability*, 17, 4894. doi:10.3390/su17114894.
- Cheng, F., Zhang, J., & Wen, C. (2016a). Cost-sensitive large margin distribution machine for classification of imbalanced data. *Pattern Recognition Letters*, 80, 107–112. <https://doi.org/10.1016/j.patrec.2016.06.009>.
- Cheng, F., Zhang, J., & Wen, C. (2016b). Cost-sensitive large margin distribution machine for classification of imbalanced data. *Pattern Recognition Letters*, 80, 107–112.
- Duan, Y. et al. (2025). Research on support vector regression short-time traffic flow prediction model for secondary roads based on associated road analysis. *Applied Sciences*, 15, 1779. doi:10.3390/app15041779.
- Guo, T., Tan, X., Yang, L., Liang, Z., Zhang, B., & Zhang, L. (2022). Domain adaptive subspace transfer model for sensor drift compensation in biologically inspired electronic nose. *Expert Systems with Applications*, 208, 118237. <https://doi.org/10.1016/j.eswa.2022.118237>.
- Gupta, U., & Gupta, D. (2021). Least squares large margin distribution machine for regression. *Applied Intelligence*, (pp. 1–36).
- Hazarika, B. B., & Gupta, D. (2023). Mode decomposition based large margin distribution machines for sediment load prediction. *Expert Systems with Applications*, (p. 120844).
- Houssein, E. H. et al. (2025). A hybrid harris hawks optimization with support vector regression for air quality forecasting. *Scientific Reports*, 15. doi:10.1038/s41598-025-86275-6.
- Hu, J., Deng, Y., & Che, J. (2025). A novel wind power interval prediction method based on neural ensemble search and dynamic conformalized quantile regression. *Applied Soft Computing*, 180, 113476. doi:10.1016/j.asoc.2025.113476.
- Hu, L., & Liu, Y. (2025). A day-ahead wind power conformal interval forecasting framework based on numerical weather prediction data and improved quantile ensemble learning. *Journal of Renewable and Sustainable Energy*, 17, 046104. doi:10.1063/5.0233770.
- Mao, Y., Song, L., & Huang, H. (2025). A novel hybrid deep learning model for day-ahead wind power interval forecasting. *Sustainability*, 17, 3239. doi:10.3390/su17073239.
- Meena, Y. et al. (2025). FLSVR: Solving lagrangian support vector regression using functional iterative method. *Neural Processing Letters*, . doi:10.1007/s11063-025-11780-8.
- Mustaffa, Z., & Sulaiman, M. H. (2025). Random forest based wind power prediction method for sustainable energy system. *Cleaner Energy Systems*, . doi:10.1016/j.cles.2025.100210.
- Rizkallah, S. (2025). Enhancing the performance of gradient boosting trees on regression problems. *Journal of Big Data*, . doi:10.1186/s40537-025-01071-3.
- Schölkopf, B., Herbrich, R., & Smola, A. J. (2001). A general representer theorem. In *Proceedings of the 14th Annual Conference on Computational Learning Theory* (pp. 416–426). Springer. doi:10.1007/3-540-44581-1_27.
- Se, H., Song, K., Liu, H., Zhang, W., Wang, X., & Liu, J. (2023). A dual drift compensation framework based on subspace learning and cross-domain adaptive extreme learning machine for gas sensors. *Knowledge-Based Systems*, 259, 110024. <https://doi.org/10.1016/j.knsys.2022.110024>.
- State Grid Renewable Energy (2022). A multi-site wind farm scada dataset for wind power forecasting. *Scientific Data*, 9, 684. doi:10.1038/s41597-022-01696-6.
- Tang, M., Ding, S. X., Yang, C., Cheng, F., Shardt, Y. A., Long, W., & Liu, D. (2019). Cost-sensitive large margin distribution machine for fault detection of wind turbines. *Cluster Computing*, 22, 7525–7537.
- Tao, J., Lai Chung, F., & Wang, S. (2012). On minimum distribution discrepancy support vector machine for domain adaptation. *Pattern Recognition*, 45, 3962–3984. <https://doi.org/10.1016/j.patcog.2012.04.014>.
- Wang, H.-z., Li, G.-q., Wang, G.-b., Peng, J.-c., Jiang, H., & Liu, Y.-t. (2019). A review of wind speed and wind power forecasting with deep neural networks. *Applied Energy*, 232, 1–17. doi:10.1016/j.apenergy.2018.09.191.
- Wang, Y., & Xu, Z. (2024). Sparse support vector quantile regression for robust estimation. *Pattern Recognition*, 147, 110089. doi:10.1016/j.patcog.2023.110089.

- Yan, J., Chen, F., Liu, T., Zhang, Y., Peng, X., Yi, D., & Duan, S. (2022). Subspace alignment based on an extreme learning machine for electronic nose drift compensation. *Knowledge-Based Systems*, 235, 107664. <https://doi.org/10.1016/j.knosys.2021.107664>.
- Yang, W., Zhang, Z., & Chen, X. (0). Vibration-based looseness identification of bolted structures via quasi-analytic wavelet packet and optimized large margin distribution machine. *Structural Health Monitoring*, 0, 14759217231159948. URL: <https://doi.org/10.1177/14759217231159948>. doi:10.1177/14759217231159948. arXiv:<https://doi.org/10.1177/14759217231159948>.
- Ye, X., Liu, C., Xiong, X., & Qi, Y. (2025a). Recurrent attention encoder-decoder network for multi-step interval wind power prediction. *Energy*, 315, 134317. doi:10.1016/j.energy.2024.134317.
- Ye, Y.-F., Wang, J., & Chen, W.-J. (2025b). Nonlinear feature selection for support vector quantile regression. *Neural Networks*, 185, 107136. doi:10.1016/j.neunet.2025.107136.
- Ye, Y.-F. et al. (2023). Twin support vector quantile regression. *Expert Systems with Applications*, 221, 121239. doi:10.1016/j.eswa.2023.121239.
- Zhai, S., Han, M., Li, Z., Yang, S., Duan, S., & Yan, J. (2023). M2fl-ccc: Multibranch multilayer feature leaning and comprehensive classification criterion for gas sensor drift compensation. *IEEE Transactions on Instrumentation and Measurement*, 72, 1–12. <https://doi.org/10.1109/TIM.2023.3296820>.
- Zhang, T., & Zhou, Z.-H. (2019). Optimal margin distribution machine. *IEEE Transactions on Knowledge and Data Engineering*, 32, 1143–1156.
- Zhang, Y., & Wang, J. (2024). Kernel density estimation based probabilistic wind power forecasting with ensemble learning. *Energy*, 289, 129912. doi:10.1016/j.energy.2023.129912.